

# Week4\_13\_ANN\_model\_adv

May 1, 2021

Artificial Neural Network (ANN)

Import Libraries

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

import warnings
warnings.filterwarnings("ignore")
```

Plot function

```
[2]: import matplotlib.pyplot as plt
plt.style.use('ggplot')

def plot_history(history):
    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']
    loss = history.history['loss']
    val_loss = history.history['val_loss']
    x = range(1, len(acc) + 1)

    plt.figure(figsize=(12, 5))
    plt.subplot(1, 2, 1)
    plt.plot(x, acc, 'b', label='Training acc')
    plt.plot(x, val_acc, 'r', label='Validation acc')
    plt.title('Training and validation accuracy')
    plt.legend()
    plt.subplot(1, 2, 2)
    plt.plot(x, loss, 'b', label='Training loss')
    plt.plot(x, val_loss, 'r', label='Validation loss')
    plt.title('Training and validation loss')
    plt.legend()
```

Import dataset

```
[4]: dataset = pd.read_csv('/Users/preethamvignesh/Downloads/ML_classwork/Week4/
↳Churn_Modelling.csv')
```

```
dataset.head()
```

```
[4]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	
3	4	15701354	Boni	699	France	Female	39	
4	5	15737888	Mitchell	850	Spain	Female	43	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

Separating Dependent and Independent variables

```
[5]: X = dataset.iloc[:, 3:13].values  
y = dataset.iloc[:, 13].values
```

Encoding categorical data

```
[6]: # Encoding categorical data  
from sklearn.preprocessing import LabelEncoder, OneHotEncoder  
from sklearn.compose import ColumnTransformer  
  
#encoding 'Geography' to number  
labelencoder_X1 = LabelEncoder()  
X[:, 1] = labelencoder_X1.fit_transform(X[:, 1])  
  
#encoding 'Gender' to numbers  
labelencoder_X2 = LabelEncoder()  
X[:, 2] = labelencoder_X2.fit_transform(X[:, 2])  
  
#encoding 'Geography' to one hot  
ct = ColumnTransformer([("Geography", OneHotEncoder(), [1])], remainder =  
    ↳ 'passthrough')  
# onehotencoder = OneHotEncoder(categorical_features = [1])  
X = ct.fit_transform(X)
```

```
# get rid of dummy variable trap
X = X[:, 1:]
```

Feature scaling

```
[7]: # Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X = sc.fit_transform(X)
```

```
[8]: from sklearn.preprocessing import MinMaxScaler
scalar = MinMaxScaler()
X= scalar.fit_transform(X)
```

Split Train and Test set

```
[9]: # Splitting the dataset into the Training set and Test set
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)
```

Model preparation with layers

```
[10]: from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Dropout

model = Sequential()
# Adding the input layer and the first hidden layer with dropout
model.add(Dense(units = 22, activation='relu', input_dim = 11)) # Input layer
    ↳and first hidden layer
# when we assing `input_dim = 11`, we actually creating input layer
# model.add(Dropout(p=0.1))

# Adding the second hidden layer with dropout
# doesn't need the input_dim params
# kernel_initializer updates weights
# activation function - rectifier
model.add(Dense(units = 22, activation='relu')) # Second hidden layer
# model.add(Dropout(p=0.1))

# Adding the output layer
# dependent variable with more than two categories (3), output_dim needs to
    ↳change (e.g. 3), activation function - softmax
model.add(Dense(1, kernel_initializer = 'glorot_uniform', activation =
    ↳'sigmoid' ))
model.add(Dense(units = 1, activation='sigmoid')) # Output layer
```

```
# Compiling the ANN - applying Stochastic Gradient Descent to whole ANN
# Several different SGD algorithms
# mathematical details based on the loss function
# binary_crossentropy, categorical_crossentropy
model.compile(loss = 'binary_crossentropy', optimizer='adam', metrics =
↳ ['accuracy'])
```

```
[11]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 22)	264
dense_1 (Dense)	(None, 22)	506
dense_2 (Dense)	(None, 1)	23
dense_3 (Dense)	(None, 1)	2

Total params: 795

Trainable params: 795

Non-trainable params: 0

Fitting Model

```
[12]: # Fitting fully connected NN to the Training set
model.fit(X_train, y_train, batch_size = 25, epochs = 100)
```

Epoch 1/100

320/320 [=====] - 1s 1ms/step - loss: 0.5992 - accuracy: 0.8032

Epoch 2/100

320/320 [=====] - 0s 1ms/step - loss: 0.5179 - accuracy: 0.8004

Epoch 3/100

320/320 [=====] - 0s 1ms/step - loss: 0.4919 - accuracy: 0.7993

Epoch 4/100

320/320 [=====] - 0s 1ms/step - loss: 0.4759 - accuracy: 0.7948

Epoch 5/100

320/320 [=====] - 0s 1ms/step - loss: 0.4527 - accuracy: 0.8041

Epoch 6/100

320/320 [=====] - 0s 1ms/step - loss: 0.4598 -

```

accuracy: 0.7911
Epoch 7/100
320/320 [=====] - 0s 1ms/step - loss: 0.4529 -
accuracy: 0.7939
Epoch 8/100
320/320 [=====] - 0s 1ms/step - loss: 0.4435 -
accuracy: 0.7956
Epoch 9/100
320/320 [=====] - 0s 1ms/step - loss: 0.4390 -
accuracy: 0.7974
Epoch 10/100
320/320 [=====] - 0s 1ms/step - loss: 0.4362 -
accuracy: 0.7959
Epoch 11/100
320/320 [=====] - 0s 1ms/step - loss: 0.4374 -
accuracy: 0.7902
Epoch 12/100
320/320 [=====] - 0s 1ms/step - loss: 0.4280 -
accuracy: 0.7964
Epoch 13/100
320/320 [=====] - 0s 1ms/step - loss: 0.4306 -
accuracy: 0.7952
Epoch 14/100
320/320 [=====] - 0s 1ms/step - loss: 0.4272 -
accuracy: 0.7872
Epoch 15/100
320/320 [=====] - 0s 1ms/step - loss: 0.4098 -
accuracy: 0.7976
Epoch 16/100
320/320 [=====] - 0s 1ms/step - loss: 0.4018 -
accuracy: 0.7969
Epoch 17/100
320/320 [=====] - 0s 1ms/step - loss: 0.3913 -
accuracy: 0.8061
Epoch 18/100
320/320 [=====] - 0s 1ms/step - loss: 0.3961 -
accuracy: 0.7984
Epoch 19/100
320/320 [=====] - 0s 1ms/step - loss: 0.3982 -
accuracy: 0.7966
Epoch 20/100
320/320 [=====] - 0s 1ms/step - loss: 0.3908 -
accuracy: 0.7988
Epoch 21/100
320/320 [=====] - 0s 1ms/step - loss: 0.3918 -
accuracy: 0.7997
Epoch 22/100
320/320 [=====] - 0s 1ms/step - loss: 0.3927 -

```

accuracy: 0.8027  
Epoch 23/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3772 -  
accuracy: 0.8296  
Epoch 24/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3925 -  
accuracy: 0.8342  
Epoch 25/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3920 -  
accuracy: 0.8400  
Epoch 26/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3903 -  
accuracy: 0.8394  
Epoch 27/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3780 -  
accuracy: 0.8472  
Epoch 28/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3775 -  
accuracy: 0.8436  
Epoch 29/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3850 -  
accuracy: 0.8407  
Epoch 30/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3776 -  
accuracy: 0.8507  
Epoch 31/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3599 -  
accuracy: 0.8541  
Epoch 32/100  
320/320 [=====] - 0s 990us/step - loss: 0.3759 -  
accuracy: 0.8456  
Epoch 33/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3721 -  
accuracy: 0.8497  
Epoch 34/100  
320/320 [=====] - 0s 988us/step - loss: 0.3622 -  
accuracy: 0.8546  
Epoch 35/100  
320/320 [=====] - 0s 992us/step - loss: 0.3669 -  
accuracy: 0.8516  
Epoch 36/100  
320/320 [=====] - 0s 993us/step - loss: 0.3654 -  
accuracy: 0.8505  
Epoch 37/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3652 -  
accuracy: 0.8545  
Epoch 38/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3710 -

```

accuracy: 0.8556
Epoch 39/100
320/320 [=====] - 0s 1ms/step - loss: 0.3528 -
accuracy: 0.8596
Epoch 40/100
320/320 [=====] - 0s 1ms/step - loss: 0.3583 -
accuracy: 0.8584
Epoch 41/100
320/320 [=====] - 0s 1ms/step - loss: 0.3665 -
accuracy: 0.8500
Epoch 42/100
320/320 [=====] - 0s 1ms/step - loss: 0.3591 -
accuracy: 0.8577
Epoch 43/100
320/320 [=====] - 0s 1ms/step - loss: 0.3595 -
accuracy: 0.8529
Epoch 44/100
320/320 [=====] - 0s 1ms/step - loss: 0.3643 -
accuracy: 0.8519
Epoch 45/100
320/320 [=====] - 0s 1ms/step - loss: 0.3679 -
accuracy: 0.8528
Epoch 46/100
320/320 [=====] - 0s 1ms/step - loss: 0.3550 -
accuracy: 0.8557
Epoch 47/100
320/320 [=====] - 0s 1ms/step - loss: 0.3583 -
accuracy: 0.8571
Epoch 48/100
320/320 [=====] - 0s 1ms/step - loss: 0.3573 -
accuracy: 0.8571
Epoch 49/100
320/320 [=====] - 0s 1ms/step - loss: 0.3625 -
accuracy: 0.8551
Epoch 50/100
320/320 [=====] - 0s 1ms/step - loss: 0.3491 -
accuracy: 0.8563
Epoch 51/100
320/320 [=====] - 0s 1ms/step - loss: 0.3481 -
accuracy: 0.8632
Epoch 52/100
320/320 [=====] - 0s 1ms/step - loss: 0.3372 -
accuracy: 0.8673
Epoch 53/100
320/320 [=====] - 0s 995us/step - loss: 0.3644 -
accuracy: 0.8500
Epoch 54/100
320/320 [=====] - 0s 1ms/step - loss: 0.3409 -

```

```

accuracy: 0.8649
Epoch 55/100
320/320 [=====] - 0s 1ms/step - loss: 0.3614 -
accuracy: 0.8500
Epoch 56/100
320/320 [=====] - 0s 1ms/step - loss: 0.3575 -
accuracy: 0.8531
Epoch 57/100
320/320 [=====] - 0s 973us/step - loss: 0.3417 -
accuracy: 0.8611
Epoch 58/100
320/320 [=====] - 0s 988us/step - loss: 0.3550 -
accuracy: 0.8566
Epoch 59/100
320/320 [=====] - 0s 1ms/step - loss: 0.3476 -
accuracy: 0.8592
Epoch 60/100
320/320 [=====] - 0s 1ms/step - loss: 0.3516 -
accuracy: 0.8585
Epoch 61/100
320/320 [=====] - 0s 1ms/step - loss: 0.3456 -
accuracy: 0.8634
Epoch 62/100
320/320 [=====] - 0s 1ms/step - loss: 0.3531 -
accuracy: 0.8620
Epoch 63/100
320/320 [=====] - 0s 973us/step - loss: 0.3429 -
accuracy: 0.8629
Epoch 64/100
320/320 [=====] - 0s 976us/step - loss: 0.3503 -
accuracy: 0.8613
Epoch 65/100
320/320 [=====] - 0s 1ms/step - loss: 0.3554 -
accuracy: 0.8550
Epoch 66/100
320/320 [=====] - 0s 1ms/step - loss: 0.3440 -
accuracy: 0.8630
Epoch 67/100
320/320 [=====] - 0s 1ms/step - loss: 0.3391 -
accuracy: 0.8669
Epoch 68/100
320/320 [=====] - 0s 1ms/step - loss: 0.3490 -
accuracy: 0.8597
Epoch 69/100
320/320 [=====] - 0s 1ms/step - loss: 0.3372 -
accuracy: 0.8642
Epoch 70/100
320/320 [=====] - 0s 1ms/step - loss: 0.3378 -

```



accuracy: 0.8663  
Epoch 71/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3364 -  
accuracy: 0.8653  
Epoch 72/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3410 -  
accuracy: 0.8629  
Epoch 73/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3388 -  
accuracy: 0.8677  
Epoch 74/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3317 -  
accuracy: 0.8703  
Epoch 75/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3468 -  
accuracy: 0.8603  
Epoch 76/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3413 -  
accuracy: 0.8591  
Epoch 77/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3429 -  
accuracy: 0.8639  
Epoch 78/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3400 -  
accuracy: 0.8582  
Epoch 79/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3330 -  
accuracy: 0.8654  
Epoch 80/100  
320/320 [=====] - 0s 996us/step - loss: 0.3348 -  
accuracy: 0.8683  
Epoch 81/100  
320/320 [=====] - 0s 1ms/step - loss: 0.3425 -  
accuracy: 0.8587  
Epoch 82/100  
320/320 [=====] - 0s 976us/step - loss: 0.3413 -  
accuracy: 0.8587  
Epoch 83/100  
320/320 [=====] - 0s 985us/step - loss: 0.3376 -  
accuracy: 0.8661  
Epoch 84/100  
320/320 [=====] - 0s 971us/step - loss: 0.3427 -  
accuracy: 0.8622  
Epoch 85/100  
320/320 [=====] - 0s 986us/step - loss: 0.3329 -  
accuracy: 0.8665  
Epoch 86/100  
320/320 [=====] - 0s 980us/step - loss: 0.3393 -

```

accuracy: 0.8616
Epoch 87/100
320/320 [=====] - 0s 1ms/step - loss: 0.3317 -
accuracy: 0.8618
Epoch 88/100
320/320 [=====] - 0s 1ms/step - loss: 0.3383 -
accuracy: 0.8601
Epoch 89/100
320/320 [=====] - 0s 980us/step - loss: 0.3366 -
accuracy: 0.8610
Epoch 90/100
320/320 [=====] - 0s 1ms/step - loss: 0.3264 -
accuracy: 0.8701
Epoch 91/100
320/320 [=====] - 0s 1ms/step - loss: 0.3351 -
accuracy: 0.8648
Epoch 92/100
320/320 [=====] - 0s 995us/step - loss: 0.3344 -
accuracy: 0.8654
Epoch 93/100
320/320 [=====] - 0s 1ms/step - loss: 0.3285 -
accuracy: 0.8711
Epoch 94/100
320/320 [=====] - 0s 979us/step - loss: 0.3390 -
accuracy: 0.8643
Epoch 95/100
320/320 [=====] - 0s 974us/step - loss: 0.3320 -
accuracy: 0.8632
Epoch 96/100
320/320 [=====] - 0s 1ms/step - loss: 0.3306 -
accuracy: 0.8680
Epoch 97/100
320/320 [=====] - 0s 982us/step - loss: 0.3366 -
accuracy: 0.8616
Epoch 98/100
320/320 [=====] - 0s 972us/step - loss: 0.3455 -
accuracy: 0.8579
Epoch 99/100
320/320 [=====] - 0s 1ms/step - loss: 0.3427 -
accuracy: 0.8569
Epoch 100/100
320/320 [=====] - 0s 990us/step - loss: 0.3331 -
accuracy: 0.8653

```

[12]: <tensorflow.python.keras.callbacks.History at 0x135b660a0>

Model Evaluate

```
[13]: model.evaluate(X_test,y_test)
```

```
63/63 [=====] - 0s 924us/step - loss: 0.3585 -  
accuracy: 0.8470
```

```
[13]: [0.3585488796234131, 0.847000002861023]
```

```
[14]: # Predicting on the Test set  
y_pred = model.predict(X_test)  
y_pred = (y_pred > 0.5)
```

```
[15]: # Get accuracy on Test set  
from sklearn.metrics import confusion_matrix  
cm_test = confusion_matrix(y_test, y_pred)  
cm_test
```

```
[15]: array([[1476,  113],  
          [ 193,  218]])
```

```
[16]: history = model.fit(X_train, y_train,  
                          epochs=20,  
                          verbose=True,  
                          validation_data=(X_test, y_test),  
                          batch_size=52)  
  
loss, accuracy = model.evaluate(X_train, y_train, verbose=False)  
print("Training Accuracy: {:.4f}".format(accuracy))  
loss, accuracy = model.evaluate(X_test, y_test, verbose=False)  
print("Testing Accuracy: {:.4f}".format(accuracy))  
plot_history(history)
```

Epoch 1/20

```
154/154 [=====] - 1s 5ms/step - loss: 0.3325 -  
accuracy: 0.8648 - val_loss: 0.3540 - val_accuracy: 0.8560
```

Epoch 2/20

```
154/154 [=====] - 0s 2ms/step - loss: 0.3312 -  
accuracy: 0.8666 - val_loss: 0.3516 - val_accuracy: 0.8500
```

Epoch 3/20

```
154/154 [=====] - 0s 2ms/step - loss: 0.3306 -  
accuracy: 0.8671 - val_loss: 0.3529 - val_accuracy: 0.8525
```

Epoch 4/20

```
154/154 [=====] - 0s 2ms/step - loss: 0.3312 -  
accuracy: 0.8655 - val_loss: 0.3519 - val_accuracy: 0.8540
```


Epoch 5/20

```
154/154 [=====] - 0s 2ms/step - loss: 0.3318 -  
accuracy: 0.8636 - val_loss: 0.3536 - val_accuracy: 0.8570
```

Epoch 6/20

```
154/154 [=====] - 0s 2ms/step - loss: 0.3315 -
```

accuracy: 0.8648 - val\_loss: 0.3523 - val\_accuracy: 0.8550  
Epoch 7/20  
154/154 [=====] - 0s 2ms/step - loss: 0.3305 -  
accuracy: 0.8670 - val\_loss: 0.3563 - val\_accuracy: 0.8425  
Epoch 8/20  
154/154 [=====] - 0s 2ms/step - loss: 0.3311 -  
accuracy: 0.8660 - val\_loss: 0.3511 - val\_accuracy: 0.8470  
Epoch 9/20  
154/154 [=====] - 0s 2ms/step - loss: 0.3317 -  
accuracy: 0.8629 - val\_loss: 0.3537 - val\_accuracy: 0.8495  
Epoch 10/20  
154/154 [=====] - 0s 2ms/step - loss: 0.3303 -  
accuracy: 0.8655 - val\_loss: 0.3520 - val\_accuracy: 0.8535  
Epoch 11/20  
154/154 [=====] - 0s 2ms/step - loss: 0.3299 -  
accuracy: 0.8664 - val\_loss: 0.3529 - val\_accuracy: 0.8545  
Epoch 12/20  
154/154 [=====] - 0s 2ms/step - loss: 0.3298 -  
accuracy: 0.8684 - val\_loss: 0.3641 - val\_accuracy: 0.8395  
Epoch 13/20  
154/154 [=====] - 0s 2ms/step - loss: 0.3312 -  
accuracy: 0.8679 - val\_loss: 0.3550 - val\_accuracy: 0.8510  
Epoch 14/20  
154/154 [=====] - 0s 2ms/step - loss: 0.3295 -  
accuracy: 0.8680 - val\_loss: 0.3506 - val\_accuracy: 0.8540  
Epoch 15/20  
154/154 [=====] - 0s 2ms/step - loss: 0.3291 -  
accuracy: 0.8658 - val\_loss: 0.3519 - val\_accuracy: 0.8525  
Epoch 16/20  
154/154 [=====] - 0s 2ms/step - loss: 0.3296 -  
accuracy: 0.8664 - val\_loss: 0.3524 - val\_accuracy: 0.8555  
Epoch 17/20  
154/154 [=====] - 0s 2ms/step - loss: 0.3294 -  
accuracy: 0.8668 - val\_loss: 0.3550 - val\_accuracy: 0.8455  
Epoch 18/20  
154/154 [=====] - 0s 2ms/step - loss: 0.3296 -  
accuracy: 0.8679 - val\_loss: 0.3525 - val\_accuracy: 0.8500  
Epoch 19/20  
154/154 [=====] - 0s 2ms/step - loss: 0.3291 -  
accuracy: 0.8679 - val\_loss: 0.3504 - val\_accuracy: 0.8525  
Epoch 20/20  
154/154 [=====] - 0s 2ms/step - loss: 0.3289 -  
accuracy: 0.8673 - val\_loss: 0.3535 - val\_accuracy: 0.8530  
Training Accuracy: 0.8705  
Testing Accuracy: 0.8530



```
[17]: y_pred = model.predict(X_test)
      y_pred = y_pred > 0.5

      con_res = confusion_matrix(y_test, y_pred)
      print("Confusion matrix:")
      print(confusion_matrix(y_test, y_pred))
      print("Accuracy: {:.2f}%".format(accuracy_score(y_test, y_pred)*100))
```


Confusion matrix:

```
[[1532  57]
 [ 237 174]]
```

Accuracy: 85.30%

Metrix: Confusion Matrix

```
[19]: %run -i '/Users/preethamvignesh/Desktop/Work/ML_EIT/
      ↪confusion_matrix_different_ways.py'
      df_confmatrx = pd.DataFrame(con_res, range(2), range(2))
      df_confmatrx
      cmap = 'PuRd'
      confusion_matrix_dfrntway(df_confmatrx, cmap=cmap, fz=17)
```



Predicting single observation

```
[18]: # Predicting a single new observation
new_prediction = model.predict(sc.transform(np.array([[0, 0, 600, 1, 40, 3,
↪ 60000, 2, 1, 1, 50000]])))
new_prediction = (new_prediction > 0.5)
print(new_prediction)
```

```
[[False]]
```

```
[ ]:
```

```
[ ]:
```