

CNN_model1

May 9, 2021

Convolutional Neural Network (CNN)

Import libraries

```
[1]: import numpy as np
import pandas as pd

import tensorflow as tf

from tensorflow.keras import datasets, layers, models
import matplotlib.pyplot as plt
```

Load the Images and split the train and test

```
[2]: (train_images, train_labels), (test_images, test_labels) = datasets.cifar10.
    ↪load_data()

# Normalize pixel values to be between 0 and 1
train_images, test_images = train_images / 255.0, test_images / 255.0
```

```
[3]: test_labels
```

```
[3]: array([[3],
           [8],
           [8],
           ...,
           [5],
           [1],
           [7]], dtype=uint8)
```

Display sample images (2x5)

```
[4]: train_images.shape
```

```
[4]: (50000, 32, 32, 3)
```

Making Labels and class names

```
[5]: class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
                    'dog', 'frog', 'horse', 'ship', 'truck']
```

```
plt.figure(figsize=(10,4))
for i in range(10):
    plt.subplot(2,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(train_images[i], cmap=plt.cm.binary)
    # The CIFAR labels happen to be arrays,
    # which is why you need the extra index
    plt.xlabel(class_names[train_labels[i][0]])
plt.show()
```



Model Architecture

- Create Convolution base simple and common pattern: CONV2D and MaxPooling2D layers.
- AS a input, CNN takes tensors of shape (image_height, image_width, and color channels)

```
[6]: model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

```
[7]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36928

Total params: 56,320
 Trainable params: 56,320
 Non-trainable params: 0

The above output of CONV2D and MaxPooling2D layer is a 3D tensor of shape(height,width, and channels).

The height and width dimensions tend to shrink as you go in the deeper in the network.

Typically, as the width and height shrink, you can afford (computationally) to add more output channels in each Conv2D layer.

ADD Dense layer on Top

To complete our model, you will feed the last output tensor from the convolutional base (of shape (4, 4, 64)) into one or more Dense layers to perform classification.

```
[8]: model.add(layers.Flatten())
      model.add(layers.Dense(64, activation='relu'))
      model.add(layers.Dense(10))
      print("Complete architecture of our model")
```

Complete architecture of our model

```
[9]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0

```

-----
conv2d_2 (Conv2D)                (None, 4, 4, 64)                36928
-----
flatten (Flatten)                (None, 1024)                    0
-----
dense (Dense)                    (None, 64)                      65600
-----
dense_1 (Dense)                  (None, 10)                      650
=====
Total params: 122,570
Trainable params: 122,570
Non-trainable params: 0
-----

```

As you can see, our (4, 4, 64) outputs were flattened into vectors of shape (1024) before going through two Dense layers.

Compile and train the model

```

[10]: model.compile(optimizer='adam',
                    loss=tf.keras.losses.
                        ↳SparseCategoricalCrossentropy(from_logits=True),
                    metrics=['accuracy'])

history = model.fit(train_images, train_labels, epochs=10,
                    validation_data=(test_images, test_labels))

```

```

Epoch 1/10
1563/1563 [=====] - 19s 12ms/step - loss: 1.7692 -
accuracy: 0.3460 - val_loss: 1.3117 - val_accuracy: 0.5303
Epoch 2/10
1563/1563 [=====] - 19s 12ms/step - loss: 1.1897 -
accuracy: 0.5773 - val_loss: 1.0436 - val_accuracy: 0.6269
Epoch 3/10
1563/1563 [=====] - 20s 13ms/step - loss: 1.0064 -
accuracy: 0.6428 - val_loss: 1.0070 - val_accuracy: 0.6443
Epoch 4/10
1563/1563 [=====] - 19s 12ms/step - loss: 0.9130 -
accuracy: 0.6792 - val_loss: 0.8990 - val_accuracy: 0.6859
Epoch 5/10
1563/1563 [=====] - 19s 12ms/step - loss: 0.8279 -
accuracy: 0.7089 - val_loss: 0.8816 - val_accuracy: 0.6940
Epoch 6/10
1563/1563 [=====] - 20s 13ms/step - loss: 0.7656 -
accuracy: 0.7315 - val_loss: 0.8858 - val_accuracy: 0.7023
Epoch 7/10
1563/1563 [=====] - 19s 12ms/step - loss: 0.7280 -
accuracy: 0.7437 - val_loss: 0.8394 - val_accuracy: 0.7118
Epoch 8/10


```

```
1563/1563 [=====] - 20s 13ms/step - loss: 0.6768 -  
accuracy: 0.7613 - val_loss: 0.9132 - val_accuracy: 0.6930  
Epoch 9/10  
1563/1563 [=====] - 19s 12ms/step - loss: 0.6291 -  
accuracy: 0.7803 - val_loss: 0.8685 - val_accuracy: 0.7119  
Epoch 10/10  
1563/1563 [=====] - 19s 12ms/step - loss: 0.5919 -  
accuracy: 0.7918 - val_loss: 0.9007 - val_accuracy: 0.7084
```

Evaluate the model

```
[11]: plt.plot(history.history['accuracy'], label='accuracy')  
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')  
plt.xlabel('Epoch')  
plt.ylabel('Accuracy')  
plt.ylim([0.5, 1])  
plt.legend(loc='lower right')  
  
test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
```

```
313/313 - 1s - loss: 0.9007 - accuracy: 0.7084
```



```
[12]: print(test_acc)
```

```
0.7084000110626221
```

```
[13]: from sklearn.metrics import accuracy_score, confusion_matrix
test_labels_pred = model.predict(test_images)
test_labels_pred.shape
```

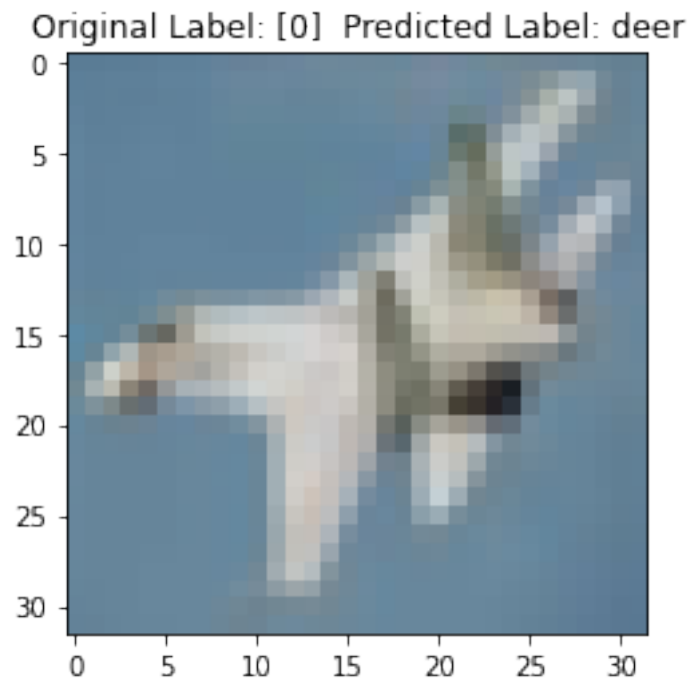
[13]: (10000, 10)

Prediction using test images


```
[14]: for i in range(10,20):
    plt.imshow(test_images[i], cmap=plt.cm.binary)

    #Image name prediction
    image = test_images[i]
    image1 = np.array(image).reshape((1,32,32,3))
    #print(image.shape, image1.shape)
    prediction = model.predict(image1)
    #print(prediction.shape)


    score = tf.nn.softmax(prediction[0])
    # print(class_names[np.argmax(score)])
    plt.title('Original Label: {0} Predicted Label: {1}'.
    ↪format(test_labels[i], class_names[np.argmax(score)]))
    plt.show()
```




Original Label: [9] Predicted Label: truck




Original Label: [5] Predicted Label: cat




Original Label: [7] Predicted Label: horse




Original Label: [9] Predicted Label: truck




Original Label: [8] Predicted Label: deer




Original Label: [5] Predicted Label: dog

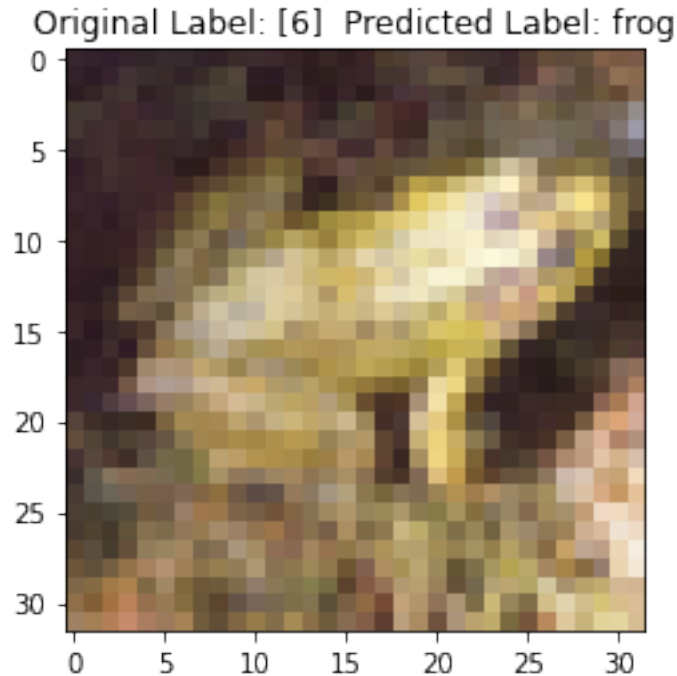


Original Label: [7] Predicted Label: horse



Original Label: [8] Predicted Label: ship






Prediction: export the image and prediction

```
[15]: from tensorflow import keras
      from tensorflow.keras import layers
      from tensorflow.keras.models import Sequential

      timg = "/home/jayanthikishore/Downloads/ML_classwork/Week5/horse.jpg"
      # timg = "/Users/preethamvignesh/Downloads/deer.jpg"
      # timg = "/Users/preethamvignesh/Downloads/dog.jpg"
      img = keras.preprocessing.image.load_img(timg, target_size=(32, 32))

      plt.imshow(img, cmap=plt.cm.binary)
      #change the dimension to train and test data (32x32)
      img_array = keras.preprocessing.image.img_to_array(img)
      img_array = tf.expand_dims(img_array, 0) # Create a batch
      #Prediction
      prediction = model.predict(img_array)
      score = tf.nn.softmax(prediction[0])
      # print(class_names[np.argmax(score)])
      plt.title('Predicted Label: {0}'.format(class_names[np.argmax(score)]))
      plt.show()
      img_array.shape
```



```
[15]: TensorShape([1, 32, 32, 3])
```

```
[ ]:
```

```
[16]: class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',  
                    'dog', 'frog', 'horse', 'ship', 'truck']
```

```
[ ]:
```